

1 Proof of Safety of Strategy for Interval Scheduling Problem

Lemma 1. *Let j be the job with the earliest finishing time. Then there is an optimum solution in which j is scheduled.*

Proof. Let S be any optimum solution for the interval selection instance. If $j \in S$, then we are done. So, we assume $j \notin S$. Let j' be the first job in S , i.e., the job with the earliest finishing time in S .

Then we show that $S' := S \setminus \{j'\} \cup \{j\}$ is also an optimum solution for the instance. First, all jobs in $S \setminus \{j'\}$ start at or after $f_{j'}$ since j' is the first job in S . By our choice of j we have $f_j \leq f_{j'}$. So, all jobs in $S \setminus \{j'\}$ start at or after f_j . Therefore, S' is a valid solution. $|S'| = |S|$ and thus S' is also valid.

In any case, there is an optimum solution that contains j . \square

2 Proof of Safety of Strategy for Offline Caching Problem

Lemma 2. *Assume at time 1 a page fault happens and there are no empty pages in the cache. Let p^* be the page in cache that is not requested until furthest in the future. Then, there is an optimum solution in which p^* is evicted at time 1.*

Proof. Let S be any optimum solution for the offline caching instance. If S evicts page p^* at time 1, then we are done. So assume S evicts some page $p' \neq p^*$ at time 1. Let t^* be the first time that p^* is requested, and t' be the first time that p' is requested; if a page is not requested in the future, then we let the correspondent time to be 0. So by our choice of p^* , we have $t' < t^*$. (An exception is that both t' and t^* are ∞ . In this case both p' and p^* are not used in the future and thus evicting p^* is equivalent to evicting p' .)

Our goal is to construct a new optimum solution S' such that

- S' evicts p^* at time 1.
- the number of page misses in S' is at most that in S .

(Notice that by the assumption that S is optimum, the second property implies that the number of page misses in S' is the same as that in S . However, in our proof, it is more convenient for us to not take the optimality of S into the account, and only show one direction described in the second property.)

For a time t , let $\text{cache}_t(S)$ and $\text{cache}_t(S')$ be the sets of pages in the solution S and S' respectively, at the end of time t . Then $|\text{cache}_t(S)| = |\text{cache}_t(S')| = k$ for any time t .

At time 1, S evicts p' and we let S' evict p^* . They both have a page fault. Moreover, $\text{cache}_1(S) \setminus \text{cache}_1(S') = \{p^*\}$ and $\text{cache}_1(S') \setminus \text{cache}_1(S) = \{p'\}$. That is, at the end of time 1, the only difference between $\text{cache}_1(S)$ and $\text{cache}_1(S')$ is that the former contains p^* and the latter contains p' .

For every time t from 2 to $t' - 1$, we copy the solution S to the solution S' . Notice that both p' and p^* are not requested in this time period. If S has a page hit, so does S' . If S has a page miss and evicts a page other than p^* , then S' has a page miss and evicts the same page. If at some time t , S evicts p^* , then we let S' evict p' , and we have $\text{cache}_t(S') = \text{cache}_t(S)$. In this case, we can copy S until the time T , to obtain the desired solution S' .

So, we assume in any time t from 2 to $t' - 1$, S does not evict the page p^* . Then, $\text{cache}_{t'-1}(S) \setminus \text{cache}_{t'-1}(S') = \{p^*\}$ and $\text{cache}_{t'-1}(S') \setminus \text{cache}_{t'-1}(S) = \{p'\}$. At time t' , S has a page miss and S' has a page hit. S evicts some page and load page p' . Thus, by time t' , S' has one less page miss than S . Moreover $|\text{cache}_{t'}(S) \cap \text{cache}_{t'}(S')| \geq k - 1$.

For every time t from $t' + 1$ to time T , we again copy S to construct our solution S' . We maintain that $|\text{cache}_t(S) \cap \text{cache}_t(S')| \geq k - 1$. Moreover, if at some time step t , S has a page hit and S' has a page miss, then we have $\text{cache}_t(S') = \text{cache}_t(S)$. Focus on time t :

- If both S and S' have a page hit, then nothing needs to be done.
- If both S and S' have a page miss, then we let S' evict the same page as S does; if S evicts the page in $\text{cache}_t(S) \setminus \text{cache}_{t'}(S')$, then let S' evict the page in $\text{cache}_t(S') \setminus \text{cache}_{t'}(S)$. We still have $|\text{cache}_t(S) \cap \text{cache}_t(S')| \geq k - 1$.

- If S has a page miss and S' has a page hit, then we do not need to do anything for S' . S will evict some page and load the page in $\text{cache}_{t-1}(S') \setminus \text{cache}_{t-1}(S)$. So, we shall have $|\text{cache}_t(S') \cap \text{cache}_t(S)| \geq k-1$. (Using the optimality of S , one can show that this case can not happen; but this is not needed.)
- If S has a page hit and S' has a page miss, then the page requested is the one in $\text{cache}_{t-1}(S) \setminus \text{cache}_{t-1}(S')$. Then, we let S' evict the page in $\text{cache}_{t-1}(S') \setminus \text{cache}_{t-1}(S)$, and load the page in $\text{cache}_{t-1}(S) \setminus \text{cache}_{t-1}(S')$. Then, we have $\text{cache}_t(S') = \text{cache}_t(S)$.

So, from time $t' + 1$ to T , S' incurs at most one more page miss than S does. Overall the number of page misses in S' is at most that in S . \square

3 Proof of Safety of Strategy for Best Prefix Code Problem

Lemma 3. *Given an instance for the best prefix code problem, and let x_1 and x_2 be the two least frequent letters. Then there is an optimum encoding tree in which the two nodes holding x_1 and x_2 are brothers.*

Proof. Let S be an optimum encoding tree for the instance. Then, there must be two deepest nodes in S that are brothers. Let v_1 and v_2 be the two nodes. If the two nodes hold x_1 and x_2 (not necessarily in this order), then we are done. Otherwise, we can make one or two swaps to make this happen.

In each swap, we take a node $u \notin \{v_1, v_2\}$ that holds some letter $x \in \{x_1, x_2\}$, and the node $v \in \{v_1, v_2\}$ that holds some letter $y \notin \{x_1, x_2\}$. Then we swap the letters held by the two nodes: we let u hold y and let v hold x . Notice that $f_x \leq f_y$ as x_1 and x_2 are the two least frequent letters, and $d_u \leq d_v$, where d_u and d_v are the depth of u and v in the solution S respectively. This holds since v_1 and v_2 are two deepest nodes in the tree. Then, $d_u f_y + d_v f_x \leq d_u f_x + d_v f_y$ since $(d_u f_x + d_v f_y) - (d_u f_y + d_v f_x) = (d_v - d_u)(f_y - f_x) \geq 0$. So, after the swap, the cost can only go down. Eventually, we obtain a new solution S' whose cost is at most that of S . Moreover, in S' , v_1 and v_2 hold x_1 and x_2 . \square